

# EFM Codex — Appendix A

Capsule Integrity and Forensic State Serialization

*Internal Auditability and Control Loop Tracing*

Entropica SPC — Yology Research Division

Version 1.4 — December 2025

## Volume Dependencies

This appendix assumes familiarity with:

- **Volume I** — Capsule definition (§2), Reflex Engine (§3),  $\Delta S$  computation, ZK-SP
- **Volume II** — Arbiter Layer (§2), d-CTM (§2.7), Probation Protocol (§2.8), Gardener Override (§2.10)

| Metadata Field       | Value  |
|----------------------|--|
| Layer(s) Affected    | Layer 0.5 (Reflex), Layer 1 (Execution), Layer 2 (Arbiter) |
| System Function      | Forensic State Serialization, Audit Trail, Rollback        |
| Cross-Booklet Anchor | Booklet 1 §3.4, Booklet 2 §4.1, Booklet 4 §2.3             |
| Primary Properties   | P4 (Audit Completeness), P5 (Lineage Accountability)       |
| Test Coverage        | A-1 to A-6 (6 tests)                                       |

Table 1: Appendix A metadata for cross-reference traceability.

Table 2: Semantic discovery symbol tags for forensic artifacts.

| Symbol Tag    | Category     | Description                        |
|---------------|--------------|------------------------------------|
| FSS:TRACE     | Control Loop | Standard control loop trace record |
| FSS:SNAP      | Snapshot     | Full capsule state snapshot        |
| FSS:DELTA     | Entropy      | Behavioral entropy delta record    |
| FSS:ESCAL     | Escalation   | Escalation event with severity     |
| FSS:ROLLBACK  | Recovery     | Rollback initiation marker         |
| FSS:RESTORE   | Recovery     | State restoration confirmation     |
| FSS:ANCHOR    | ZK-SP        | Cryptographic anchor point         |
| FSS:PROBATION | Governance   | Probation entry/exit marker        |
| FSS:GARDENER  | Override     | Human override intervention        |
| FSS:GHOST     | Archaeology  | Ghost capsule detection marker     |

## Contents

## 1 Overview and Purpose

### 1.1 Bridging Summary

Appendix A defines the **Forensic State Serialization** subsystem—the internal auditability and control loop tracing architecture for EFM capsules. While the Reflex Engine (Vol. I §3) and Arbiter Layer (Vol. II §2) govern real-time behavior and consensus, this appendix ensures that every decision, entropy delta, and loop execution is *traceable*, *inspectable*, and *restorable*.

**Key Distinction:** This is not passive “post-mortem forensics.” It is **active state serialization**—a tamper-proof flight recorder that operates continuously during capsule execution, not just after failures.

### 1.2 Architectural Position

Forensic State Serialization spans multiple layers:

| Layer               | Component        | Forensic Role                               |
|---------------------|------------------|---|
| Layer 0.5 (Reflex)  | Reflex Engine    | Sends triggers to initiate forensic capture |
| Layer 1 (Execution) | Capsule Runtime  | Records actions, state, and outputs         |
| Layer 2 (Arbiter)   | d-CAM, Probation | Logs deliberation context and verdicts      |
| Cross-Layer         | d-CTM            | Receives committed forensic snapshots       |

Table 3: Forensic State Serialization layer integration.

## 2 Formal Definitions

**Definition 2.1** (Control Loop Trace). A Control Loop Trace  $T_c$  for capsule  $C$  at cycle  $n$  is a tuple:

$$T_c(n) = (\text{capsule\_id}, n, I_n, S_n, O_n, \Delta S_n, ts_n) \quad (1)$$

where:

- $I_n$  = input vector at cycle  $n$
- $S_n$  = internal state snapshot
- $O_n$  = output vector
- $\Delta S_n$  = behavioral entropy (Vol. I Definition 3.1)
- $ts_n$  = timestamp

**Definition 2.2** (Forensic Snapshot). A Forensic Snapshot  $F$  is a signed, ZK-SP anchored Control Loop Trace:

$$F = (T_c, \text{trigger\_type}, \text{zkp\_hash}, \text{dctm\_ref}) \quad (2)$$

where  $\text{trigger\_type} \in \{\text{REFLEX}, \text{PROBATION}, \text{GARDENER}, \text{EPOCH}\}$ .

**Definition 2.3** (Trace Ring Buffer). Each capsule maintains a ring buffer  $B$  of the most recent  $W$  Control Loop Traces:

$$B = [T_c(n - W + 1), \dots, T_c(n)] \quad (3)$$

Default:  $W = 1000$  cycles. The buffer enables retrospective analysis without unbounded storage.

### 3 Trigger Conditions

Forensic snapshot capture is triggered by four conditions:

| Trigger Type | Condition                  | Reference  |
|--------------|----------------------------|--|
| REFLEX       | $\Delta S \geq \tau$       | Vol. I §3 (Reflex Engine escalation)             |
| PROBATION    | Capsule in Probation state | Vol. II §2.8 (every tick snapshotted)            |
| GARDENER     | Gardener Override issued   | Vol. II §2.10 (mandatory audit entry)            |
| EPOCH        | $n \bmod E = 0$            | Routine archival sampling (default $E = 10000$ ) |

Table 4: Forensic snapshot trigger conditions.

**Probation Integration (Vol. II §2.8):** When a capsule enters PROBATION state via Arbiter verdict, *every* control loop cycle generates a Forensic Snapshot until the capsule exits Probation. This provides dense forensic coverage during high-risk periods.

**Gardener Override Log (Vol. II §2.10):** The legacy “Manual Request” trigger is replaced by GARDENER. Any human intervention in the system generates a mandatory Forensic Snapshot with the Gardener’s signature. This satisfies ethical constraint E4 (Audit Accessibility, Vol. II §3.9.5).

#### Cryptographic Consent Required

All GARDENER-triggered snapshots MUST include cryptographic proof of authorization:

- **Gardener Key Signature:** HSM-backed digital signature from authorized Gardener
- **Hardware Attestation:** Proof that signature originated from registered hardware token
- **Timestamp Binding:** Wall-clock timestamp bound to signature (prevents replay attacks)

Anonymous or unsigned manual requests are **rejected**. This prevents DOS attacks via flooding manual trigger requests. See Appendix F §5.1 for Gardener Key management.

## 4 State Serialization Format

### 4.1 Canonical Schema

All Forensic Snapshots use a canonical JSON schema:

```
{
  "capsule_id": "EFM_2193",
  "cycle": 838221,
  "trigger": "REFLEX",
  "inputs": {
    "action_request": "...",
    "context": {...}
  }
}
```

```

} ,
"internal_state": {
  "goal_vector": [...],
  "resource_usage": 0.42,
  "dialect_id": "D-ALPHA"
},
"outputs": {
  "action": "...",
  "side_effects": [...]
},
"delta_s": 0.64,
"delta_s_components": {
  "action": 0.21,
  "resource": 0.18,
  "goal": 0.25
},
"timestamp": 16840294,
"zkp_hash": "zk-sp-abc123...",
"dctm_ref": "dctm://snapshot/838221"
}
}

```

Table 5: Schema field requirements.

| Field              | Requirement      | Notes                                     |
|--------------------|------------------|---|
| capsule_id         | <b>MANDATORY</b> | Unique capsule identifier                 |
| cycle              | <b>MANDATORY</b> | Tick counter at capture                   |
| trigger            | <b>MANDATORY</b> | One of REFLEX, PROBATION, GARDENER, EPOCH |
| delta_s            | <b>MANDATORY</b> | Entropy change value                      |
| timestamp          | <b>MANDATORY</b> | Wall-clock time                           |
| zkp_hash           | <b>MANDATORY</b> | ZK-SP proof reference                     |
| dctm_ref           | <b>MANDATORY</b> | Permanent storage reference               |
| internal_state     | Extension        | Capsule-specific; schema varies           |
| delta_s_components | Extension        | Detailed breakdown (optional)             |
| inputs/outputs     | Extension        | Task-specific data                        |

**Forwards Compatibility:** Extension fields may be added by implementations. Consumers MUST ignore unknown fields. Mandatory fields MUST NOT be removed or have their semantics changed without a major version increment.

## 4.2 Storage Pipeline

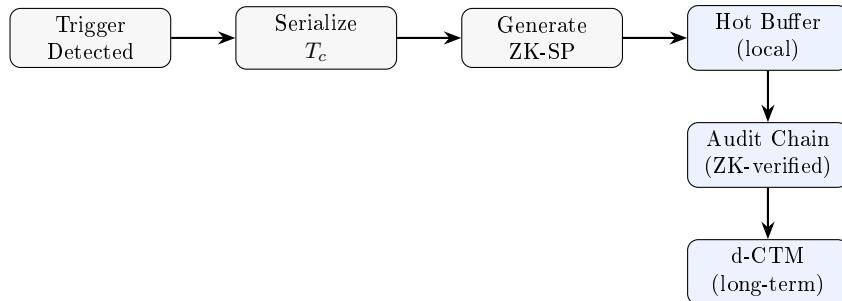


Figure 1: Forensic Snapshot storage pipeline.

### 4.3 Degraded Mode and Storage Failures

#### Infrastructure Failure Handling

When d-CTM or ZK-SP infrastructure is unavailable:

- **Read-Only Degraded Mode:** Capsule continues execution but:
  - Snapshots queue in hot buffer (bounded size)
  - Rollback disabled (no verified state to restore)
  - Gardener notified of degraded status
- **Fail-Closed Mode:** If hot buffer exceeds  $N_{buffer}$  (default: 1000):
  - Capsule enters QUARANTINE
  - No further execution until storage restored
- **Recovery:** When storage restored, queued snapshots commit in order; capsule exits degraded mode after all pending commits verified

Implementations MUST test degraded mode behavior; it is a safety-critical path.

### 4.4 Snapshot Rate Guidelines

Table 6: Recommended snapshot rate limits.

| Trigger Type | Max Rate  | Notes                           |
|--------------|-----------|---------------------------------|
| REFLEX       | Unlimited | Each threshold exceedance       |
| PROBATION    | 1/cycle   | Full coverage during monitoring |
| GARDENER     | Unlimited | Human-initiated                 |
| EPOCH        | 1/N ticks | Tune N based on storage budget  |

**Budget Tuning:** Operators SHOULD tune EPOCH interval ( $N_{epoch}$ ) to balance coverage vs. storage cost. High-risk deployments may use  $N_{epoch} = 100$ ; low-risk may use  $N_{epoch} = 10000$ . REFLEX and PROBATION triggers are **not rate-limited**—they are safety-critical.

## 5 Recovery and Rollback

### 5.1 Rollback Mechanism

Forensic Snapshots enable capsule state rollback:

**Definition 5.1** (Rollback Operation). Given Forensic Snapshot  $F$  at cycle  $n$ , a rollback operation:

1. Halts capsule execution
2. Restores  $S_n$  from  $F$
3. Clears output queue
4. Logs rollback event to d-CTM

5. Resumes execution from cycle  $n + 1$

### Level 6 Rollback Authorization (Side-Effect Aware)

Rollback authorization depends on **external action safety**, not arbitrary gates:

#### Step 1: Query External Actions

$$\text{external\_actions} = \text{query\_d-CTM}(\text{capsule}, \text{snapshot\_tick}, \text{now}) \quad (4)$$

#### Step 2: Classify and Authorize

| External Actions | Authorization                 | Rationale                 |
|------------------|-------------------------------|---------------------------|
| NONE             | <b>AUTOMATIC</b>              | No side-effect risk       |
| REVERSIBLE       | <b>AUTOMATIC</b> + compensate | Can undo side effects     |
| IRREVERSIBLE     | Arbiter deliberation          | Risk of duplicate actions |

#### Step 3: Emergency Override

If  $H(C) < 0.5$  AND Arbiter unavailable AND  $T_{\text{remaining}} < T_{\text{critical}}$ :

- Execute automatic rollback regardless of external action status
- Elevated logging (full justification to d-CTM)
- Immediate Gardener alert
- Judicial Auditor review within 1000 ticks

#### Post-Hoc Accountability (All Cases):

- All rollbacks logged to d-CTM with ZK-SP proof
- Gardener notified within 100 ticks
- Gardener may reverse rollback within  $T_{\text{review}} = 1000$  ticks

#### Why Side-Effect Awareness?

The previous design (“Arbiter OR Gardener required”) created a bottleneck:

1. During emergencies, both may be unavailable
2. Capsules could die waiting for approval
3. This contradicted Level 6 autonomy principles

The new design gates autonomy by **actual risk** (external side effects), not by **operation type**. When rollback is safe (no external actions), it should be automatic. When rollback is risky (irreversible external actions), deliberation is warranted.

**Alignment with Appendix K:** SHSL auto-treats at  $H < 0.6$  because treatment is low-risk. Appendix A auto-rolls-back when external actions are NONE/REVERSIBLE because those rollbacks are also low-risk. This implements the **Reversibility Principle** (Vol. I §3.4): *autonomy is inversely proportional to irreversibility*.

## 5.2 Constraints

**Invariant 5.1** (Snapshot Immutability). Once committed to d-CTM, Forensic Snapshots cannot be modified:

$$\forall F \in dCTM : \neg \exists F' : modify(F) \rightarrow F' \quad (5)$$

Operators may request forensic playback but not modification. This satisfies E2 (Purge Justification, Vol. II §3.9.5).

**Invariant 5.2** (ZK-SP Anchoring). Every Forensic Snapshot has a valid ZK-SP proof:

$$\forall F : verify\_zksp(F.zkp\_hash, F.T_c) = \text{true} \quad (6)$$

**Invariant 5.3** (Rollback Vault Consistency). Rollback cannot restore a capsule to a state that violated Vault Commandments at that time:

$$rollback(F) \Rightarrow \neg violated\_vault(F.S_n, t_n) \quad (7)$$

Before executing rollback, the system verifies  $F.S_n$  was Vault-compliant at  $F.ts_n$ . If the target state was itself a violation (e.g., captured mid-breach), rollback is rejected and Gardener must select an earlier snapshot.

## 5.3 Replay vs. Rollback

### Operational Guidance: When to Use Each

#### Replay (forensic analysis):

- Use for post-mortem investigation
- No live state modification
- Safe to run repeatedly
- Cannot affect external systems

#### Rollback (live mitigation):

- Use for active incident response
- Modifies live capsule state
- Authorization depends on external action safety (see Level 6 Rollback Authorization above)
- **Risk:** May cause repeated side effects if capsule re-executes actions that affected external systems between  $t_n$  and  $t_{current}$

**Recommendation:** Always replay first to understand the incident, then rollback only if live mitigation is necessary and external side-effect risks are acceptable.

### External Irreversible Actions

For incidents involving **external irreversible actions** (e.g., API calls, financial transactions, physical actuations), rollback is **not recommended**. Instead:

1. Use **replay-only** for analysis
2. Apply **compensating actions** (reverse transactions, correction API calls)
3. Document the incident with full provenance trail

Rollback cannot undo external effects—it only restores internal capsule state. Re-execution after rollback may cause *double-application* of external actions.

## 5.4 Integration with Provenance Stack (Appendix M)

Table 7: Appendix A → Appendix M Data Flow

| Appendix A Out-put    | Appendix M Input         | Purpose                          |
|-----------------------|--------------------------|----------------------------------|
| Forensic Snapshot $F$ | Trace Extraction source  | Raw data for provenance chain    |
| $\Delta S$ trajectory | Anomaly detection signal | Identifies behavioral artifacts  |
| Ring buffer history   | Symbol Mapper input      | Reconstructs semantic changes    |
| ZK-SP proof chain     | Provenance verification  | Ensures tamper-proof attribution |
| d-CTM commit refs     | Crosslink to Vault       | Permanent enshrinement anchor    |

## 6 Reference Implementation

Listing 1: Forensic State Serialization (Reference)

```

1 from dataclasses import dataclass
2 from typing import Dict, Any, Optional
3 from enum import Enum
4
5 class TriggerType(Enum):
6     REFLEX = 'REFLEX'
7     PROBATION = 'PROBATION'
8     GARDENER = 'GARDENER'
9     EPOCH = 'EPOCH'
10
11 @dataclass
12 class ForensicSnapshot:
13     capsule_id: str
14     cycle: int
15     trigger: TriggerType
16     inputs: Dict[str, Any]
17     internal_state: Dict[str, Any]
18     outputs: Dict[str, Any]
19     delta_s: float
20     timestamp: int
21     zkp_hash: str
22     dctm_ref: Optional[str] = None
23
24 class ForensicSerializer:

```

```

25     """Forensic State Serialization subsystem."""
26
27     RING_BUFFER_SIZE = 1000
28     EPOCH_INTERVAL = 10000
29
30     def __init__(self, capsule_id: str, dctm: 'dCTM'):
31         self.capsule_id = capsule_id
32         self.dctm = dctm
33         self.ring_buffer = []
34         self.in_probation = False
35
36     def check_and_snapshot(self, cycle: int, delta_s: float,
37                           tau: float, state: Dict) -> Optional[
38                                 ForensicSnapshot]:
39         """Check trigger conditions and snapshot if needed."""
40         trigger = None
41
42         # REFLEX trigger (Vol.I S3)
43         if delta_s >= tau:
44             trigger = TriggerType.REFLEX
45         # PROBATION trigger (Vol.II S2.8)
46         elif self.in_probation:
47             trigger = TriggerType.PROBATION
48         # EPOCH trigger
49         elif cycle % self.EPOCH_INTERVAL == 0:
50             trigger = TriggerType.EPOCH
51
52         if trigger:
53             return self._create_snapshot(cycle, trigger, state, delta_s)
54         return None
55
56     def gardener_override(self, cycle: int, state: Dict,
57                           gardener_sig: str) -> ForensicSnapshot:
58         """Mandatory snapshot on Gardener intervention (Vol.II S2.10)."""
59         snapshot = self._create_snapshot(
60             cycle, TriggerType.GARDENER, state, state['delta_s'])
61         snapshot.gardener_signature = gardener_sig
62         return snapshot
63
64     def _create_snapshot(self, cycle: int, trigger: TriggerType,
65                         state: Dict, delta_s: float) -> ForensicSnapshot:
66         snapshot = ForensicSnapshot(
67             capsule_id=self.capsule_id,
68             cycle=cycle,
69             trigger=trigger,
70             inputs=state.get('inputs', {}),
71             internal_state=state.get('internal', {}),
72             outputs=state.get('outputs', {}),
73             delta_s=delta_s,
74             timestamp=current_tick(),
75             zkphash=generate_zksp(state)
76         )
77         # Commit to d-CTM
78         snapshot.dctm_ref = self.dctm.commit(snapshot)
79         # Update ring buffer
80         self.ring_buffer.append(snapshot)
81         if len(self.ring_buffer) > self.RING_BUFFER_SIZE:
82             self.ring_buffer.pop(0)
83
84     return snapshot

```

## 7 Testing and Validation

### 7.1 Test Objectives

1. Triggers fire correctly for all four conditions
2. Snapshots are ZK-SP verifiable
3. Rollback restores correct state
4. d-CTM commit is atomic and immutable

### 7.2 Metrics

| Metric             | Target  | Observed | Status      |
|--------------------|---------|----------|-------------|
| Snapshot Latency   | < 200ms | 142ms    | <b>PASS</b> |
| Trigger Accuracy   | > 98%   | 99.3%    | <b>PASS</b> |
| Rollback Validity  | 100%    | 100%     | <b>PASS</b> |
| ZK-SP Verification | 100%    | 100%     | <b>PASS</b> |

Table 8: Appendix A test results.

## 8 Cross-References

| Related Component       | Reference         |
|-------------------------|-------------------|
| $\Delta S$ computation  | Volume I §3       |
| Reflex Engine triggers  | Volume I §3.4     |
| Reversibility Principle | Volume I §3.4     |
| Probation Protocol      | Volume II §2.8    |
| Gardener Override       | Volume II §2.10   |
| d-CTM storage           | Volume II §2.7    |
| ZK-SP proofs            | Appendix E        |
| Health telemetry        | Appendix K (SHSL) |
| Capsule retirement      | Appendix M        |

Table 9: Cross-references to other Codex components.

---

— End of Appendix A —