

# EFM Codex: Operator's Guide

Deployment, Monitoring, and Response

*Appendix O to Volumes I-II*

**Entropica SPC**

Yology Research Division

Version 1.1 — December 2025

**Audience:** This guide is for operators, SREs, and governance teams responsible for deploying and running an EFM forest. It is *not* for proving invariants or modifying the specification.

**Scope:** Summarizes normative behavior from Volumes I-II and translates it into operational concepts and procedures. For formal definitions and proofs, see the source volumes.

## Contents

<b>1 Core Concepts in Operator Terms</b>	<b>2</b>
1.1 Entity Glossary . . . . .	2
1.2 The Two Primary Health Signals . . . . .	2
<b>2 Standing Up an EFM Forest</b>	<b>3</b>
2.1 Pre-Deployment Checklist . . . . .	3
2.2 First Forest Verification . . . . .	3
<b>3 Day-to-Day Monitoring</b>	<b>4</b>
3.1 Key Metrics and Alerts . . . . .	4
3.2 Dashboard Layout . . . . .	4
<b>4 Response Playbooks</b>	<b>5</b>
4.1 Entropy Spike (S2–S3 Severity) . . . . .	5
4.2 SCI Below Threshold . . . . .	5
4.3 Orphan Detection . . . . .	6
4.4 Catastrophic Failure: Arbiter Unavailability . . . . .	6
4.5 Catastrophic Failure: d-CTM Compromise . . . . .	7
4.6 Catastrophic Failure: Runaway Spawn . . . . .	7
4.7 Cascading Failure Decision Tree (Gap 7) . . . . .	8
4.8 Partial System Failure . . . . .	10
<b>5 Configuration and Tuning</b>	<b>11</b>
5.1 Tunable Parameters . . . . .	11
5.2 Safe Change Workflow . . . . .	11

<b>6</b>	<b>Governance and Audits</b>	<b>12</b>
6.1	Gardener Responsibilities . . . . .	12
6.2	Incident Reconstruction . . . . .	12
6.3	Compliance Checklist . . . . .	12
<b>7</b>	<b>Quick Reference</b>	<b>13</b>
7.1	Common Commands . . . . .	13
7.2	Severity Levels . . . . .	13
<b>8</b>	<b>Implementation Guidance</b>	<b>13</b>
8.1	Micro-Heuristic Deployment . . . . .	13
8.2	ZK-SP Audit Path Verification . . . . .	14
8.3	Deployment Hooks . . . . .	14

# 1 Core Concepts in Operator Terms

## 1.1 Entity Glossary

EFM Term	Operator Analog	Description
Capsule	Managed agent instance	Individual autonomous unit with role and permissions
Vault (Layer 0)	System-wide safety baseline	Immutable constraints; cannot be modified at runtime
Reflex Engine (L0.5)	Per-capsule circuit breaker	Sub-10ms safety interrupt; blocks unsafe actions
Arbiter Layer (L2)	Distributed safety council	Deliberative consensus for ambiguous decisions
d-CAM	Consensus protocol	Byzantine-fault-tolerant voting among Arbiters
Forest (L3)	Fleet topology	Swarm structure with dialect lineages
Trunk	Main branch	Primary dialect shared by capsule group
Branch	Canary/variant	Diverged dialect subset under evaluation
Gardener	Human operator	Override authority with cryptographic keys

Table 1: EFM concepts mapped to operator-familiar terms.

## 1.2 The Two Primary Health Signals

Operators monitor two key metrics:

1.  **$\Delta S$  (Behavioral Entropy)** — Per-capsule (Vol. I §3). Measures how “strange” an individual capsule is behaving relative to its baseline. Threshold:  $\tau$  (role-dependent, typically 0.3–0.5).
2. **SCI (Swarm Coherence Index)** — Per-trunk/branch (Vol. II §3.2). Measures how “aligned” a population is behaving collectively. Range: 0–1. Threshold:  $\theta_{SCI}$  (deployment-dependent, typically 0.65–0.85).

**Key Insight:** A swarm can have all capsules with  $\Delta S < \tau$  (all individually “safe”) while  $SCI < \theta_{SCI}$  (collectively drifting). Both signals are necessary.

## 2 Standing Up an EFM Forest

### 2.1 Pre-Deployment Checklist

1. **Define Vault Commandments** — Immutable safety rules. Once deployed, cannot be changed without full system rebuild.
2. **Configure Genesis Parameters** — Role archetypes,  $\tau/\lambda$  defaults, permission matrices.
3. **Provision Capsules** — Deploy initial capsule fleet with Genesis Protocol. Verify each capsule's Vault Hash Chain links to system root.
4. **Deploy Arbiter Set** — Minimum 5 Judicial Capsules for  $f = 2$  Byzantine tolerance. Configure d-CAM quorum parameters.
5. **Initialize d-CTM** — Distributed ledger storage. Verify append-only property and ZK-SP integration.
6. **Configure Gardener Keys** — Generate and securely store Gardener override keys. Test override on non-critical verdict.
7. **Enable SCI Computation** — Configure  $d_{behav}$  components, sampling strategy (if  $> 10^4$  capsules), and alert thresholds.

### 2.2 First Forest Verification

Before production traffic:

Check	Expected	Command/Method
Capsules online	$\geq N_{min}$	Dashboard: Active Capsules
d-CTM reachable	All nodes	<code>dctm-health -all</code>
SCI computation	Value $\in [0, 1]$	Dashboard: SCI gauge
Reflex latency	$< 10\text{ms}$ p99	Metrics: <code>reflex_latency_p99</code>
Arbiter quorum	$\geq 5$ available	<code>dcam-status -quorum</code>
Gardener override	Success	Manual test on staging verdict

Table 2: First-forest verification checklist.

## 3 Day-to-Day Monitoring

### 3.1 Key Metrics and Alerts

Wire up the following to your monitoring stack (Prometheus, Datadog, etc.):

Metric	Alert Threshold	Action
<code>delta_s_p95</code>	$> \tau - 0.05$	Review high-entropy capsules
<code>sci_per_trunk</code>	$< \theta_{alert}$	Investigate drift pattern
<code>reflex_halt_rate</code>	$> 10/\text{min}$ sustained	Check for misconfiguration
<code>arbiter_verdict_latency</code>	$> 5\text{s p99}$	Scale Arbiter pool
<code>orphan_count</code>	$> 0$	Verify Orphan Protocol active
<code>fork_rate</code>	$> 2/\text{hour}$	Possible instability
<code>branch_count</code>	$> B_{max} - 2$	Approaching branch limit

Table 3: Recommended monitoring configuration.

### 3.2 Dashboard Layout

Recommended panels:

1. **Fleet Health** — Capsule count by state (ACTIVE, PROBATION, QUARANTINE, PURGED)
2. **Entropy Distribution** — Histogram of  $\Delta S$  across all capsules
3. **SCI Timeseries** — Per-trunk SCI with threshold lines
4. **Arbiter Activity** — Verdict rate, latency, decision breakdown
5. **Forest Topology** — Visual tree of trunks/branches with capsule counts
6. **Gardener Actions** — Recent overrides and approvals

## 4 Response Playbooks

### 4.1 Entropy Spike (S2–S3 Severity)

#### Playbook: High $\Delta S$ on Individual Capsule

**Trigger:**  $\Delta S \geq \tau$  for capsule  $C$ , Reflex halt issued.

**Steps:**

1. **Inspect History:** Query d-CTM for  $C$ 's recent action log. Identify which  $\Delta S$  component spiked (action, resource, goal).
2. **Classify:**
  - *Expected workload change* → Adjust  $\tau$  via Threshold Governance (Vol. II §2.6)
  - *Legitimate anomaly* → Leave halt in place, monitor
  - *Misconfiguration* → Fix configuration, restart capsule
  - *Potential compromise* → Escalate to Arbiter for QUARANTINE
3. **Document:** Log decision rationale in d-CTM via Gardener annotation.

**Escalation:** If pattern repeats across multiple capsules, investigate swarm-level cause (see SCI Dip playbook).

### 4.2 SCI Below Threshold

#### Playbook: SCI Dip Below $\theta_{fork}$

**Trigger:**  $SCI < \theta_{fork}$  for  $> T_{debounce}$  ticks.

**Steps:**

1. **Identify Divergence:** Use DDI breakdown to find which capsule subset is drifting. Check behavioral distance matrix.
2. **Review Arbiter Verdicts:** Query recent FORK proposals. If Arbiter already deliberating, wait for verdict.
3. **Validate Fork Decision:**
  - If FORK approved: Verify branch created correctly, SCI improved per-group
  - If FORK rejected: Review DCG rationale; consider manual Gardener review
4. **Post-Fork Monitoring:** Track branch stability. Prepare for eventual Merge if branch proves valuable.

**Escalation:** If SCI continues dropping post-fork, investigate deeper (possible adversarial drift or systemic issue).

### 4.3 Orphan Detection

#### Playbook: Orphan Capsules Detected

**Trigger:** `orphan_count > 0`.

**Steps:**

1. **Verify Protocol Active:** Confirm Orphan Protocol is running (capsules should be in QUARANTINE).
2. **Check Recovery Attempts:** Query d-CTM for re-sync attempts. Are compatible trunks being tried?
3. **Monitor Timeout:** If approaching  $N_{orphan\_timeout}$ , decide:
  - *Recoverable* → Manually assign to trunk via Gardener action
  - *Unrecoverable* → Allow PURGE, review forensics
4. **Post-Purge Review:** Examine d-CTM entry for PURGE. Verify compliance with E2 Purge Justification constraint (Vol. II §3.9.5).

**Escalation:** Multiple orphans simultaneously may indicate trunk corruption or network partition.

### 4.4 Catastrophic Failure: Arbiter Unavailability

#### Runbook 1: ALL Arbiters Unavailable

**Trigger:** Arbiter availability < 20% or all Arbiters unreachable.

**Severity:** S4 (CRITICAL)

**Immediate Actions:**

1. **Verify Scope:** Confirm this is not a monitoring failure. Check network connectivity to Arbiter nodes.
2. **Emergency Spawn:** If truly unavailable, spawn emergency Arbiters from standby pool:

```
gardener spawn-arbiters --emergency --count=3
```

3. **Queue Freeze:** Escalation queue will grow. Set queue priority to CRITICAL-only:

```
efm-cli escalation --mode=critical-only
```

4. **Capsule Guidance:** Broadcast to capsules: “Arbiters recovering, reduce escalations”

**Recovery Criteria:**

- Arbiter availability > 50%
- Escalation queue < 100 pending

- No new Arbiter failures for 10K ticks

**Post-Incident:** Root cause analysis required. Check for common cause (resource exhaustion, coordinated attack, network partition).

## 4.5 Catastrophic Failure: d-CTM Compromise

### Runbook 2: d-CTM Integrity Compromised

**Trigger:** d-CTM hash verification failure OR write latency > 1000ms sustained.

**Severity:** S5 (EMERGENCY)

#### Immediate Actions:

1. **Halt All Operations:** This is a Layer 0 violation. Issue forest-wide FREEZE:

```
gardener forest --action=FREEZE --reason=DCTM_COMPROMISE
```

2. **Isolate d-CTM:** Prevent further writes until integrity verified:

```
efm-cli dctm --mode=read-only
```

3. **Integrity Check:** Run full verification against backup:

```
efm-cli dctm --verify --against=backup
```

4. **If Corrupted:** Initiate rollback to last known-good state (see Appendix J §8).

#### Recovery Criteria:

- d-CTM integrity verified (hash match)
- Write latency < 100ms
- Gardener sign-off on recovery

**Post-Incident:** Mandatory forensic replay (Appendix M). Identify attack vector or corruption source.

## 4.6 Catastrophic Failure: Runaway Spawn

### Runbook 3: Runaway Spawn Event

**Trigger:** Spawn rate  $> 2 \times R_{max}$  OR depth exceeds  $D_{max}$  anywhere.

**Severity:** S4 (CRITICAL)

#### Immediate Actions:

1. **Emergency Lockdown:** ASG should auto-trigger, but verify:

```
efm-cli spawn --check-lockdown
```

2. **If Not Locked:** Manually trigger:

```
gardener spawn --lockdown --R_max=0 --tau=1.0
```

3. **Identify Source:** Query d-CTM for spawn lineage:

```
efm-cli lineage --depth-exceeded --last=1000
```

4. **Quarantine Branch:** Isolate the runaway lineage:

```
gardener quarantine --lineage=<ROOT_ID>
```

#### Recovery Criteria:

- Spawn rate  $\leq R_{max}$
- All capsules  $D \leq D_{max}$
- ASG parameters restored to operational bounds

**Post-Incident:** Review ASG configuration. Was  $R_{max}$  appropriate? Adjust deployment profile if needed.

## 4.7 Cascading Failure Decision Tree (Gap 7)

### Critical Operational Guidance

Real-world failures cascade. When multiple symptoms appear simultaneously, use this decision tree to identify the **primary failure** and execute the correct runbook.

#### Runbook 0: Cascading Failure Diagnosis

**Trigger:** Multiple failure symptoms detected simultaneously (e.g., Arbiter failures + escalation queue growth + d-CTM slowdown).

##### The Problem:

Tick 0: Arbiter failure -> Escalation queue grows  
 Tick 1K: Queue growth -> Memory pressure -> d-CTM writes slow  
 Tick 2K: d-CTM slow -> Health monitoring delayed -> False readings  
 Tick 3K: False health -> ASG reduces spawn -> System underutilized  
 Tick 4K: Multiple symptoms: Which runbook?

##### Step 1: Identify Primary Failure

Run diagnostics in this **exact priority order**:

1. **d-CTM Integrity Check** (highest priority):

```
efm-cli dctm --verify-integrity
```

*If CORRUPTED → Runbook 2 IMMEDIATELY* (this is Layer 0)

## 2. Reflex Engine Health:

```
efm-cli reflex --status
```

*If DOWN → EMERGENCY HALT* (Layer 0 violation)

## 3. Arbiter Availability:

```
efm-cli arbiters --status
```

*If < 50% → Runbook 1* (Arbiter recovery)

## 4. Spawn Rate Anomaly:

```
efm-cli spawn --check
```

*If > R<sub>max</sub> → Runbook 3* (spawn lockdown)

### Step 2: Execute Primary Runbook

Follow the procedures for whichever failure was identified first. Do NOT try to address secondary effects until primary is resolved.

### Step 3: Monitor Secondary Effects

Every 1000 ticks, re-run Step 1 diagnostics:

- If all checks pass → Recovery in progress
- If new failure detected → Escalate to Gardener (human decision needed)
- If same failure persists → Continue current runbook

### Step 4: Root Cause Analysis

After recovery, mandatory forensic replay (Appendix M):

- Which component failed first?
- How did failure propagate?
- How to prevent cascade in future?

### Decision Matrix:

Symptom Combination	Likely Primary	Action
Arbiter ↓ + Queue ↑	Arbiter failure	Runbook 1
d-CTM slow + Everything else	d-CTM issue	Runbook 2
Spawn ↑ + Depth ↑	Runaway spawn	Runbook 3
All symptoms at once	d-CTM first	Runbook 2 then reassess
Reflex down + Anything	Reflex failure	EMERGENCY HALT

## 4.8 Partial System Failure

### Playbook: Partial Arbiter Failure (20–50%)

**Trigger:** Arbiter availability between 20% and 50%.

**Severity:** S3 (HIGH)

**Note:** This is NOT “all Arbiters failed” (Runbook 1) but partial degradation.

**Steps:**

#### 1. Assess Severity:

- > 50% available → Minor degradation, monitor only
- 20 – 50% available → This playbook
- < 20% available → Runbook 1 (emergency)

#### 2. Spawn Replacement Arbiters:

```
gardener spawn-arbiters --count=<FAILED_COUNT/2>
```

#### 3. Redistribute Load:

```
efm-cli arbiters --rebalance
```

#### 4. Investigate Common Cause:

- Network partition? Check connectivity between Arbiter nodes
- Resource exhaustion? Check memory/CPU on failed nodes
- Coordinated attack? Check for suspicious patterns in d-CTM

#### 5. Alert Gardener: Partial failures often indicate emerging problems

### Recovery Criteria:

- Arbiter availability > 90%
- No new failures for 5K ticks
- Escalation queue stable

## 5 Configuration and Tuning

### 5.1 Tunable Parameters

Parameter	Default	Tuning Method	Constraint
$\tau$ (entropy threshold)	Role-dependent	Threshold Governance	$\tau \in [\tau_{min}, \tau_{max}]$
$\lambda$ (heuristic tolerance)	0.1	Threshold Governance	$\lambda \in [0, \lambda_{max}]$
$\theta_{SCI}$ (fork threshold)	0.7	Operator config	Per deployment regime
$T_{debounce}$	100 ticks	Operator config	$> 0$
$N_{orphan\_timeout}$	5000 ticks	Operator config	$> T_{recovery}$
$B_{max}$ (branch limit)	10	Operator config	$> 0$
ATP horizon $K_{max}$	5000 ticks	Operator config	$> 0$

Table 4: Tunable parameters and constraints.

### 5.2 Safe Change Workflow

All parameter changes follow this workflow:

1. **Propose:** Submit change request with justification
2. **Validate:** For  $\tau/\lambda$  changes, run Monotonic Sensitivity check (Vol. I Theorem 3.1)
3. **Approve:** Gardener signs change (or Arbiter for Threshold Governance)
4. **Log:** Change recorded to d-CTM with ZK-SP proof
5. **Monitor:** Watch affected metrics for  $T_{observation}$  ticks
6. **Rollback:** If regression detected, revert via same workflow

#### Never Modify Directly

The following are **immutable** and cannot be changed at runtime:

- Vault Commandments (Layer 0)
- Reflex-Core rules (Layer 0.5)
- Genesis root hash
- d-CTM append-only property

Changes to these require full system rebuild.

## 6 Governance and Audits

### 6.1 Gardener Responsibilities

1. **Override Authority:** Can override any d-CAM verdict (subject to Vault constraints)
2. **DCG Review:** Review Deliberation Context Graphs for high-impact decisions (Vol. II §2.3)
3. **Micro-Heuristic Approval:** Sign off on Reflex-Heuristic updates derived from Arbiter precedent
4. **Conflict Arbitration:** Resolve disputes between safety and performance
5. **Audit Cooperation:** Provide d-CTM access for compliance reviews

### 6.2 Incident Reconstruction

To reconstruct any incident from audit trail:

1. **Identify Verdict:** Query d-CTM by capsule ID, timestamp, or verdict type
2. **Retrieve DCG:** Use `dcg_hash` to fetch full reasoning graph (Vol. II Definition 2.3)
3. **Verify ZK-SP:** Confirm proof chain is valid (tamper detection)
4. **Trace Lineage:** Follow Vault Hash Chain to genesis for affected capsules
5. **Document:** Produce audit report showing invariant compliance

### 6.3 Compliance Checklist

For “meaningful human control” (EU AI Act Article 14 and similar):

Requirement	EFM Mechanism	Evidence
Human can understand system behavior	DCG, d-CTM logs	Audit reports
Human can intervene	Gardener override	Override logs in d-CTM
Human can stop system	Gardener PURGE authority	Purge records
Decisions are traceable	ZK-SP chain	Proof verification
System cannot override human	Vault constraint on overrides	Architecture

Table 5: Regulatory compliance mapping.

## 7 Quick Reference

### 7.1 Common Commands

Task	Command
Check capsule status	<code>efm-ctl capsule status &lt;id&gt;</code>
Query d-CTM	<code>dctm-query -filter "capsule=&lt;id&gt;"</code>
View SCI for trunk	<code>efm-ctl forest sci &lt;trunk_id&gt;</code>
List active branches	<code>efm-ctl forest branches</code>
Initiate Gardener override	<code>gardener-ctl override &lt;verdict_id&gt; -reason "..."</code>
Check Arbiter quorum	<code>dcam-status -quorum</code>
Export audit trail	<code>dctm-export -from &lt;ts&gt; -to &lt;ts&gt; -format json</code>

Table 6: Common operator commands (implementation-specific; adjust for your deployment).

### 7.2 Severity Levels

Level	$\Delta S$ Range	Action	Response Time
S0 (Normal)	$< 0.5\tau$	None	—
S1 (Elevated)	$0.5\tau - 0.8\tau$	Monitor	Next shift
S2 (Warning)	$0.8\tau - \tau$	Investigate	1 hour
S3 (Critical)	$\geq \tau$	Reflex halt, review	Immediate

Table 7: Severity levels and response expectations.

## 8 Implementation Guidance

This section provides working examples for key operational concepts.

### 8.1 Micro-Heuristic Deployment

Micro-Heuristics are compact behavioral patterns deployed to the Reflex-Heuristic layer:

[language=Python, caption=Micro-Heuristic deployment example.] from efm.reflex import MicroHeuristic, ReflexHeuristic

```
Define a new micro-heuristic for resource hoarding detection hoarding_detector = MicroHeuristic(name = "resource_hoarding_v1", pattern = b"1234...56", 64 - bytesignatureseverity = Severity.S2, response = Response.THROTTLE, metadata = "author": "discovery_stack", "source": "golden_heuristic_GH - 2847", "validated": True)
```

```
Deploy via Validation Gate (ensures Monotonic Sensitivity) reflex = ReflexHeuristic.current() result = reflex.deploy(hoarding_detector)
```

```
if result.success: print(f"Deployed: {result.signature_id}") print(f"New sensitivity: {result.new_sensitivity}") else: print(f"Rejected: {result.reason}") e.g., "monotonicity violation"
```

**Monotonic Sensitivity:** The Validation Gate ensures that new micro-heuristics only *increase* detection sensitivity. If a proposed heuristic would reduce sensitivity (fewer detections), it is rejected. See Vol. I §3.5.

## 8.2 ZK-SP Audit Path Verification

Operators can verify the ZK-SP audit trail for any capsule:

```
[language=Python, caption=ZK-SP audit path verification.] from efm.audit import ZK-SPVerifier, AuditPath
```

```
Retrieve audit path for a capsule capsule_id = "C - 2024 - 12 - 15 - 00042" path = AuditPath.for_capsule(capsule_id)
```

```
Verify the complete chain verifier = ZKSPVerifier() result = verifier.verify_chain(path)
print(f"Chain length: {len(path.proofs)}") print(f"Genesis anchor: {path.genesis_hash[:16]}...") print(f"Current anchor: {path.current_hash[:16]}...") print(f"Verified: {result.valid}")
if not result.valid: print(f"Break at proof {result.break_index}") print(f"Reason: {result.reason}")
```

```
Verify specific state at tick tick = 50000 state_proof = verifier.verify_state_at_tick(capsule_id, tick) print(f"State at tick {tick} = {state_proof.state_hash[:16]}...")
```

Table 8: ZK-SP verification CLI commands.

Command	Description
efm-cli zksp verify -capsule=<id>	Verify capsule audit chain
efm-cli zksp anchor -tick=<n>	Get anchor hash at tick
efm-cli zksp replay -from=<t1> -to=<t2>	Forensic replay
efm-cli zksp export -capsule=<id>	Export proof chain

## 8.3 Deployment Hooks

Pre-deployment validation hooks ensure system integrity:

```
[language=Python, caption=Deployment hook configuration.] from efm.deploy import DeploymentHooks, ValidationResult
```

```
class ProductionHooks(DeploymentHooks):
    """Custom hooks for production deployment."""

```

```
    def pre_genesis(self, config) -> ValidationResult:
        """Validate before Genesis ceremony."""
        """Check all spawn conditions are configurable if not config.spawn_conditions"""
        return ValidationResult.FAIL("Missing spawn conditions")

    Verify ZK-SP infrastructure if not self.zksp_available():
    return ValidationResult.FAIL("ZK-SP not initialized")
    return ValidationResult.PASS

    def post_genesis(self, root_capsule) -> ValidationResult:
        """Validate after Genesis."""
        """Verify Vault binding if not root_capsule.vault_bound()"""
        return ValidationResult.FAIL("Vault binding failed")

    Verify d-CTM initialized if not root_capsule.dctm_active():
    return ValidationResult.FAIL("d-CTM not active")
    return ValidationResult.PASS

    def pre_spawn(self, parent, child_spec) -> ValidationResult:
        """Validate before any spawn."""
        """Delegate to standard spawning gate return self.spawn_gate.validate(parent, child_spec)"""
        Register hooks deploy = Deployment(profile=Profile.PRODUCTION)
        deploy.register_hooks(ProductionHooks())

```

## Changelog

v1.2 (December 2025) — Catastrophic Failure Runbooks

- **Runbook 1:** Arbiter unavailability procedures (§4.4)
- **Runbook 2:** d-CTM compromise response (§4.5)
- **Runbook 3:** Runaway spawn event handling (§4.6)
- **Gap 7 (Added):** Cascading failure decision tree (§4.7)
- **Added:** Partial system failure playbook (§4.8)
- 5 new operational playbooks total

## v1.0 (December 2025)

- Initial operator's guide
- Core concepts and glossary
- Day-to-day monitoring procedures
- Response playbooks (Entropy Spike, SCI Dip, Orphan Detection)
- Configuration and tuning guidance
- Governance and audit requirements

---

— *End of Operator's Guide* —

For formal definitions and proofs, see Volumes I-II.

For Constitutional Kernel details, see Appendix J.

For Adaptive Spawn Governance, see Appendix N.